

Concurrency Bugs (Ch 32)

Discussion Questions

- What is an atomicity-violation bug?
 - non-atomic memory updates
 - ATM example
 - how do we fix them in general?
- What is an order-violation bug?
 - a specific ordering of operations is required, but not enforced
 - how is this generally fixed?
 - What are the four conditions that must hold for deadlock to occur? (explain in your own words)
- What are the differences in attitudes of deadlock prevention, avoidance, and detect/recover?
- How can we prevent circular waiting?
 - problems?

- How can we prevent hold-and-wait?
 - grab all locks at once, atomically
 - what are the problems with this approach?
- How can we prevent no preemption?
 - try lock & release all other locks if cannot obtain
 - how is this similar to hold-and-wait prevention technique?
- How can we prevent mutual exclusion?
 - wait-free data structures
 - is this a practical approach?
- What is livelock?
- Which of the four deadlock conditions does deadlock avoidance via scheduling attempt to remove?
 - Why is this impractical?
- Which approach do you think is the most common?

- bury head in sand and hope they deadlock

doesn't happen

- What is the primary difference in the causes of non-deadlock bugs vs. deadlock bugs?

- non-deadlocking => not enough control (locks, CV, etc...)
- deadlock => using control wrong